

Machine Learning Engineer Nanodegree

Capstone Project

Chandana Neerukonda
October 17 2017

Domain Background:

Autonomous Vehicles are one of the most in vogue today. Autonomous Driving touches many challenges in different fields, especially computer vision and machine learning. One of the most talked is pedestrian detection.

Pedestrian detection systems may be possible to implement without the use of spinning light detection and ranging (LIDAR) devices, which can often cost more than \$10,000. Instead, cheap video cameras may be used for the same purpose of recognizing and tracking nearby pedestrians.[1] Research Papers show impressive results using transfer learning but the most common errors for detectors are as follows:

- Detecting tree leaves or traffic lights in background as pedestrian
- Detecting the same person twice
- Not detecting small persons

Other approaches like [2] use a real-time algorithm for the detection and tracking of image regions that possibly contain pedestrians. It then uses a classification algorithm based on the typical motion patterns of a pedestrian's legs. These approaches assume the visibility of the legs and works for walking people only!

There has been a lot of work done on Caltech Dataset[3] which is considerably the largest Pedestrian dataset.

Problem statement:

The problem is Pedestrian Detection on Caltech data set. I am personally motivated to explore the implementation of mobile nets for pedestrian detection to have this feature on an Android mobile. I believe this problem is solvable because of the new advances including SSD, YOLO and Faster RCNN and the support from Tensor flow for object detection => TFDetect (<https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/android>)

Datasets and inputs:

Caltech Dataset (http://www.vision.caltech.edu/Image_Datasets/CaltechPedestrians/)

The training data (set00-set05) consists of six training sets (~1GB each), each with 6-13 one-minute long seq files, along with all annotation information . The testing data (set06-set10) consists of five sets, again ~1GB each. Annotations for the entire dataset are also provided.

Solution Statement

I aim to extract the images from the .seq files and the corresponding annotations from the .vbb files provided. I am planning to build tf_record files for both training and testing and use them to train my model using the pre-trained object_detection module on the coco_dataset.

Benchmark Model

http://www.vision.caltech.edu/Image_Datasets/CaltechPedestrians/files/PAMI12pedestrians.pdf

This paper does benchmarking of several algorithms on the Caltech dataset.

But my approach would be entirely different from the bench_mark model because I am planning to use SSD_MobileNet / Faster_RCNN for training.


Evaluation Metrics:

The Evaluation metric I would propose would be the mAP value (detection and classification accuracy).

But the reference paper gave the evaluation metrics in terms of missrate. I do know how exactly to compare these two metrics.

Data Exploration:

I plotted the ground truth bounding boxes on the images to do a sanity-check on the bounding box parameters.



I see the existence of multiple boxes because I mixed the labels for both person and people classes.

The following table is taken from Piotr Dollar et. al. 's paper on the state of art pedestrian detectors which indicates that the CalTech data set is the largest and the most diverse pedestrian dataset available by the time the paper got published.

		imaging setup	Training			Testing			Height			Properties						
			# pedestrians	# neg. images	# pos. images	# pedestrians	# neg. images	# pos. images	10% quantile	median	90% quantile	color images	per-image eval.	no select. bias	video seqs.	temporal corr.	occlusion labels	publication
MIT	16	photo	924	-	-	-	-	-	128	128	128	✓						2000
USC-A	17	photo	-	-	-	313	-	205	70	98	133		✓					2005
USC-B	17	surv.	-	-	-	271	-	54	63	90	126		✓					2005
USC-C	18	photo	-	-	-	232	-	100	74	108	145		✓					2007
CVC	19	mobile	1000	6175 [†]	-	-	-	-	46	83	164	✓		✓				2007
TUD-det	20	mobile	400	-	400	311	-	250	133	218	278	✓	✓					2008
Daimler-CB	21	mobile	2.4k	15k [†]	-	1.6k	10k [†]	-	36	36	36			✓				2006
NICTA	22	mobile	18.7k	5.2k	-	6.9k	50k [†]	-	72	72	72	✓		✓				2008
INRIA	23	photo	1208	1218	614	566	453	288	139	279	456	✓						2005
ETH	24	mobile	2388	-	499	12k	-	1804	50	90	189	✓	✓	✓	✓			2007
TUD-Brussels	25	mobile	1776	218	1092	1498	-	508	40	66	112	✓	✓	✓				2009
Daimler-DB	26	mobile	15.6k	6.7k	-	56.5k	-	21.8k	21	47	84		✓		✓	✓		2009
Caltech	27	mobile	192k	61k	67k	155k	56k	65k	27	48	97	✓	✓	✓	✓	✓	✓	2009

Like I mentioned in the Solution_statement, I am planning to extract the images from the .seq files and the corresponding annotations from the .vbb files provided. I am planning to build tf_record files for both training and testing and use them to train my model using the pre-trained object_detection module on the coco_dataset.

Data Pre-Processing and Implementation:

The ground_truth annotations have different classes like Person, People, person(?) for the detection with lesser confidence and also occlusion_information. I am planning to not worry about the occluded detections but to consider them as true_positives. I am also planning to eliminate the person(?) class annotations and to merge the person and people detections. So, basically I will have just one class person. I might have false_negatives in the ground_truth data. But I would like to go with it just to simplify the process.

Techniques:

An seq file is a series of concatenated image frames with a fixed size header. Matlab routines for reading/writing/manipulating seq files can be found in Piotr's Matlab Toolbox (version 3.20 or later required). These routines can also be used to extract an seq file to a directory of images.

I have extracted all the images from all the sets (set00-set10) using the matlab toolbox and the annotations from the .vbb files. But ,the annotations obtained are the text files for each image and the format of the annotation is as follows:

```
% bbGt version=3people 54 149 68 95 0 0 0 0 0 0person 259 160 22 66 0 0 0 0 0 0person 18
177 11 31 0 0 0 0 0 0person 321 112 235 369 0 0 0 0 0 0
```

People and person being the classes and the the 4 numbers after the class are the (l,t,w,h) left, top, width and height respectively.

The other digits indicated presence of an occlusion and if yes, give the (l,t,w,h) values of the visible part of the person/people. The primary difference of this video bounding box annotation from the static annotation is that each object can exist for multiple frames i.e. the vbb annotation not only provides the locations of objects but also tracking information.

Since, I am not considering occlusion information, I wrote a python script to make a csv file by reading each text annotation file for each image.

The CSV file has the annotations in the following format :

filename	width	height	class	xmin	ymin	xmax	ymax
set04_V005_I00025.jp	640	480	person	112	196	131	243
g							

Now, 267994 detections are selected randomly for training and 66977 detections for testing. Also notice that since the images are randomly selected, the images are no longer temporally related to each other. Now, I go through these CSV files and read through the “filename” column to grab the corresponding images from the pool of images and place them in the corresponding train_images and test_images folders.

Algorithms:

SSD Object detection

Creating TFRecords - Tensorflow Object Detection API:

From [4], I create TFRecords files using the [generate_tfrecord.py](#) available from the raccoon detection example [5]. The only modification that I had to make here was in the class_text_to_int function. I changed this to the class “person”.

Implementation:

Next, I set up the object detection API by cloning the tensorflow object detection API (models) on my computer. (<https://github.com/tensorflow/models.git>). I installed the object detection library and generate the tf_records for both training data and test data.

Instead of training a model from scratch, I have decided to use transfer learning by using a pre-trained model on coco-object detection dataset.

TensorFlow has quite a few pre-trained models with checkpoint files available, along with configuration files. The object API also provides some [sample configurations](#) to choose from. I chose `ssd_mobilenet_v1_pets.config` file, and have changed the following parameters:

```
num_classes: 1
initial_learning_rate: 0.0004
```

And gave the corresponding input path for the `train.record` file. I also had to define the `object_detection` label text file as follows:


```
item {
  id: 1
  name: 'person'
}
```

And now I start executing the `train.py` file that comes with the package.

```
INFO:tensorflow:Global step 8869: loss = 6.8710 (9.660 sec/step)
INFO:tensorflow:Global step 8870: loss = 5.9467 (10.984 sec/step)
INFO:tensorflow:Recording summary at step 8870.
INFO:tensorflow:Global step 8871: loss = 5.2865 (9.622 sec/step)
INFO:tensorflow:Global step 8872: loss = 6.2749 (8.947 sec/step)
INFO:tensorflow:Global step 8873: loss = 6.8044 (8.757 sec/step)
INFO:tensorflow:Global step 8874: loss = 6.5693 (8.723 sec/step)
INFO:tensorflow:Global step 8875: loss = 6.9764 (9.369 sec/step)
INFO:tensorflow:Global step 8876: loss = 6.2448 (9.997 sec/step)
INFO:tensorflow:Global step 8877: loss = 5.6655 (9.480 sec/step)
INFO:tensorflow:Global step 8878: loss = 6.9724 (8.779 sec/step)
INFO:tensorflow:Global step 8879: loss = 5.9309 (8.896 sec/step)
INFO:tensorflow:Global step 8880: loss = 5.7498 (8.754 sec/step)
INFO:tensorflow:Global step 8881: loss = 6.3677 (8.640 sec/step)
INFO:tensorflow:Global step 8882: loss = 5.9362 (8.812 sec/step)
INFO:tensorflow:Global step 8883: loss = 4.7120 (9.384 sec/step)
INFO:tensorflow:Recording summary at step 8883.
INFO:tensorflow:Global step 8884: loss = 6.6195 (10.298 sec/step)
INFO:tensorflow:Global step 8885: loss = 6.1897 (8.623 sec/step)
INFO:tensorflow:Global step 8886: loss = 6.0748 (8.972 sec/step)
INFO:tensorflow:Global step 8887: loss = 6.4351 (8.774 sec/step)
INFO:tensorflow:Global step 8888: loss = 5.5326 (9.571 sec/step)
INFO:tensorflow:Global step 8889: loss = 6.2666 (9.779 sec/step)
INFO:tensorflow:Global step 8890: loss = 5.3877 (9.460 sec/step)
INFO:tensorflow:Global step 8891: loss = 5.7830 (9.904 sec/step)
INFO:tensorflow:Global step 8892: loss = 6.7693 (9.588 sec/step)
INFO:tensorflow:Global step 8893: loss = 5.9825 (9.276 sec/step)
INFO:tensorflow:Global step 8894: loss = 6.6749 (8.816 sec/step)
INFO:tensorflow:Global step 8895: loss = 6.6698 (9.168 sec/step)
INFO:tensorflow:Recording summary at step 8895.
INFO:tensorflow:Global step 8896: loss = 6.1680 (11.944 sec/step)
INFO:tensorflow:Global step 8897: loss = 5.6896 (9.424 sec/step)
INFO:tensorflow:Global step 8898: loss = 6.3189 (9.187 sec/step)
INFO:tensorflow:Global step 8899: loss = 5.5403 (8.613 sec/step)
INFO:tensorflow:Global step 8900: loss = 6.4032 (8.648 sec/step)
INFO:tensorflow:Global step 8901: loss = 5.5857 (8.902 sec/step)
INFO:tensorflow:Global step 8902: loss = 5.5389 (8.641 sec/step)
INFO:tensorflow:Global step 8903: loss = 5.5748 (8.696 sec/step)
INFO:tensorflow:Global step 8904: loss = 5.9367 (8.903 sec/step)
INFO:tensorflow:Global step 8905: loss = 5.3821 (8.748 sec/step)
INFO:tensorflow:Global step 8906: loss = 5.9318 (8.783 sec/step)
INFO:tensorflow:Global step 8907: loss = 6.1828 (8.786 sec/step)
INFO:tensorflow:Global step 8908: loss = 5.9977 (8.955 sec/step)
INFO:tensorflow:Global step 8909: loss = 6.2039 (9.088 sec/step)
INFO:tensorflow:Recording summary at step 8909.
INFO:tensorflow:Global step 8910: loss = 6.1644 (10.629 sec/step)
INFO:tensorflow:Global step 8911: loss = 6.4961 (9.044 sec/step)
INFO:tensorflow:Global step 8912: loss = 6.1769 (8.933 sec/step)
INFO:tensorflow:Global step 8913: loss = 7.0327 (8.795 sec/step)
INFO:tensorflow:Global step 8914: loss = 7.4127 (8.899 sec/step)
INFO:tensorflow:Global step 8915: loss = 6.2296 (9.375 sec/step)
INFO:tensorflow:Global step 8916: loss = 5.4407 (8.649 sec/step)
INFO:tensorflow:Global step 8917: loss = 7.1497 (9.224 sec/step)
INFO:tensorflow:Global step 8918: loss = 6.8064 (9.389 sec/step)
INFO:tensorflow:Global step 8919: loss = 6.0975 (9.333 sec/step)
INFO:tensorflow:Global step 8920: loss = 6.1756 (8.943 sec/step)
INFO:tensorflow:Global step 8921: loss = 5.2171 (9.512 sec/step)
INFO:tensorflow:Saving checkpoint to path training/model.cpkt
INFO:tensorflow:Global step 8922: loss = 6.0655 (9.616 sec/step)
INFO:tensorflow:Recording summary at step 8922.
INFO:tensorflow:Global step 8923: loss = 6.7517 (11.321 sec/step)
```

I left the model to train over night and I see that the total loss almost plateaued at around 5. I think ideally, the total loss reaching under 2 or even 1 would be great.

TensorBoard analysis on training :



So, I have generated a frozen inference graph and followed the demo object_detection jupyter notebook to load the frozen tensor flow model into memory and have loaded the label map.


I have tested on the couple of images and the detection results are as follows:






In the above , the first two images had really good detection results. In the second image, the biker was detected twice.

The eval.py script provided by tensor flow object detection package helps evaluate the model and also provides the evaluation checkpoints. I used the tensorboard at a different port to see the results.




Negative example:

The following pedestrian with the dog was not detected by the model.



The mAP@0.5IOU Precision graph is as follows: The average value is around 0.27

Precision/mAP@0.5IOU



References:

- [1] <http://www.govtech.com/fs/Googles-New-Pedestrian-Recognition-Tech-Could-Quicken-Self-Driving-Cars-to-Market.html>
- [2] <https://pdfs.semanticscholar.org/b406/6e7b3c227bcf8549af551edb163cedc1b931.pdf>
- [3] http://www.vision.caltech.edu/Image_Datasets/CaltechPedestrians/files/algorithms.pdf
- [4] <https://pythonprogramming.net/creating-tfrecord-files-tensorflow-object-detection-api-tutorial/?completed=/custom-objects-tracking-tensorflow-object-detection-api-tutorial/>
- [5] <https://towardsdatascience.com/how-to-train-your-own-object-detector-with-tensorflows-object-detector-api-bec72ecfe1d9>